

CSC3002F Networking Assignment 2

Packet Sniffing with Wireshark

1 Introduction

This assignment focuses on the 'packets' being sent around on a computer network. You will discover what is sent around on the transmission medium and what those packets 'look like'.

You may have heard of IP packets (datagrams), but do you know how it looks like and what exactly is 'inside' one? Or when sending an HTTP request to open a web page, how is that put across the transmission medium? The Wireshark tool, a packet 'sniffer', will be used to obtain answers to such questions.

The assignment is based on two '[Wireshark labs](#)' developed by Jim Kurose and Ken Ross as an adjunct to their book "Computer Networking: A Top Down Approach". Part one is the TCP lab and part two the IP lab. (An introduction to Wireshark is available on the Vula page for this assignment.)

The assignment is partly automatically marked and partly manually marked. Automatically marked questions/parts are indicated with the suffix '[AM]' while manually marked parts are indicated with the suffix [MM]

You have to submit, in ONE zip file:

1. A trace file called '`tcptrace.pcap`' containing your trace for the first half of the assignment (questions 1-9) ;
2. A trace file called '`iptrace.pcap`' containing your trace for the second half of the assignment (questions 10-18) ;
3. A structured text file called '`automark.txt`' containing answers to automatically marked questions. (File format below; sample on Vula assignment page.)
4. A document (Word/Open Office/pdf from LaTeX/...) containing answers (optionally including annotated screenshots) to manually marked questions.
You should name it '`manualmark.xxx`' where the filename extension depends on the document format.

NOTE

- You can submit **only TEN times** and the automatic marker will **only indicate that an answer is right or wrong**, so be careful and think and check twice, or even thrice, before submitting!
- Your traces must be saved as '`pcap`' files, NOT '`pcapng`' files.

The *format* for automark.txt is as follows:

```
# q1
<Trace frame number where you found the answer.>
<Source IP>
<Source port>

# q2
<Trace frame number where you found the answer.>
<Destination IP>
<Destination port>

#q3
<Trace frame number where you found the answer.>
<Sequence number>

#q4
<Trace frame number.>
<Sequence number>
<Acknowledgement number>

#q5
<Trace frame number>
<Sequence number>

#q6
<Trace frame number for 1st segment>
<Time first segment sent>
<Time ACK received>
<RTT>
<Second segment RTT>
<Estimated RTT>

#q7
<Length of first>
<Length of second>
<Length of third>
<length of fourth>

#q8
<Minimum advertised buffer space>

#q10
<Trace frame number>
<IP Address>

# q11
<TTL value>

# q12
<Protocol value>

# q13
<header length>
<body length>

# q14
<Fragmented>
```

Continued

A *sample* file looks like this:

```
# Question One
271
197.239.129.141
59536

# Question two
271
128.119.245.12
80

# Question three
213
2841915405

# Question four
220
2223153991
2841915406

# Question five
222
2841915406

# Question six
222
7.879
8.098
0.249
0.249
0.249

# Question seven
634
1400
1400
2780

# Question eight
15872

# Question 10
9
10.0.2.15

# Question 11
1

# Question 12
17
# Question 13
20
36

# Question 14
no
```


Continued

The file may contain blank line and single line comments – a comment is preceded by a hash, '#', symbol. Real numbers should be recorded to three decimal places.

Wireshark Lab:

TCP v8.1¹

Supplement to *Computer Networking: A Top-Down Approach*, 8th ed., J.F. Kurose and K.W. Ross



"Tell me and I forget. Show me and I remember. Involve me and I understand." Chinese proverb

© 2005-2021, J.F Kurose and K.W. Ross, All Rights Reserved

In this lab, we'll investigate the behaviour of the TCP protocol in detail. We'll do so by analysing a trace of the TCP segments sent and received in transferring a text file from your computer to a remote server. We'll study TCP's use of sequence and acknowledgement numbers for providing reliable data transfer; we'll see TCP's congestion control algorithm – slow start and congestion avoidance – in action; and we'll look at TCP's receiver-advertised flow control mechanism. We'll also briefly consider TCP connection setup, and we'll investigate the performance (throughput and round-trip time) of the TCP connection between your computer and the server.

Before beginning this lab, you'll probably want to review sections 3.5 and 3.7 in the text². There is also a separate Wireshark introduction pdf on Vula that you may want to read through. It provides a short overview of the tool. (There's also the manual online.)

Since the Wireshark developers keep developing the software, and there may be slight differences across OSs, and it may be that the screenshots below aren't exactly the same as what you'll see.

¹ With modifications for UCT Networks assignment 2.

² References to figures and sections are for the 8th edition of our text, *Computer Networks, A Top-down Approach*, 8th ed., J.F. Kurose and K.W. Ross, Addison-Wesley/Pearson, 2020. Our website for this book is http://gaia.cs.umass.edu/kurose_ross You'll find lots of interesting open material there.

Continued

1. Capturing a bulk TCP transfer from your computer to a remote server

Before beginning our exploration of TCP, we'll need to use Wireshark to obtain a packet trace of the TCP transfer of a file from your computer to a remote server. You'll do so by accessing a Web page that will allow you to enter the name of a file stored on your computer (which contains a project Gutenberg ASCII text), and then transfer the file to a Web server using the HTTP POST method (see section 2.2.3 in the text). We're using the POST method rather than the GET method as we'd like to transfer a large amount of data *from* your computer to another computer. Of course, we'll be running Wireshark during this time to obtain the trace of the TCP segments sent and received from your computer.

Do the following:

- Start up your web browser. Go to Vula and retrieve your assigned Gutenberg book in txt format. Store this as a .txt file somewhere on your computer.
- Next, go to <http://gaia.cs.umass.edu/wireshark-labs/TCP-wireshark-file1.html>.
- You should see a screen that looks like Figure 1.

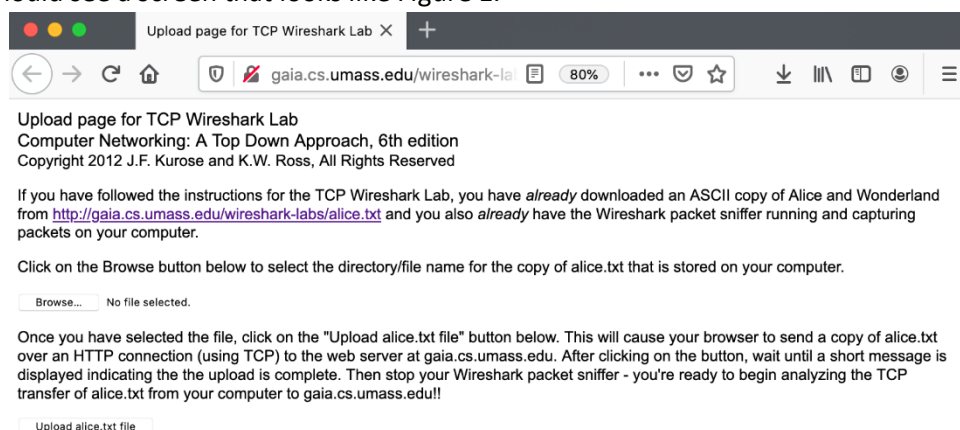



Figure 1: Page to upload the alice.txt file from your computer to gaia.cs.umass.edu

- Use the *Browse* button in this form to the file on your computer that you just created. Don't press the *"Upload alice.txt file"* button yet.³
- Now start up Wireshark and begin packet capture (see the earlier Wireshark labs if you need a refresher on how to do this).
- Returning to your browser, press the *"Upload alice.txt file"* button to upload the file to the gaia.cs.umass.edu server. Once the file has been uploaded, a short congratulations message will be displayed in your browser window.
- Stop Wireshark packet capture. Your Wireshark window should look similar to the window shown in Figure 2.

³ The button presumes that you're uploading the Gutenberg text for Alice in Wonderland; however, we have modified the lab requiring you to upload a personally assigned text.

Continued



Congratulations!

You've now transferred a copy of `alice.txt` from your computer to `gaia.cs.umass.edu`. You should now stop Wireshark packet capture. It's time to start analyzing the captured Wireshark packets!

Figure 2: Success! You've uploaded a file to `gaia.cs.umass.edu` and have hopefully captured a Wireshark packet trace while doing so.

Note: You can download a packet trace that was captured while following the steps above on one of the author's computers⁴. In addition, to your own trace, you may find it helpful when you explore the questions below.

2. A first look at the captured trace

Before analysing the behaviour of the TCP connection in detail, let's take a high-level view of the trace.

Let's start by looking at the HTTP POST message that uploaded the `alice.txt` file to `gaia.cs.umass.edu` from your computer. Find that file in your Wireshark trace, and expand the HTTP message so we can take a look at the HTTP POST message more carefully. Your Wireshark screen should look something like Figure 3.

⁴ You can download the zip file <http://gaia.cs.umass.edu/wireshark-labs/wireshark-traces-8.1.zip> and extract the trace file `tcp-wireshark-trace1-1`. This trace file can be used to answer this Wireshark lab without actually capturing packets on your own. This trace was made using Wireshark running on one of the author's computers, while performing the steps indicated in this Wireshark lab. Once you've downloaded a trace file, you can load it into Wireshark and view the trace using the *File* pull down menu, choosing *Open*, and then selecting the trace file name.

Continued

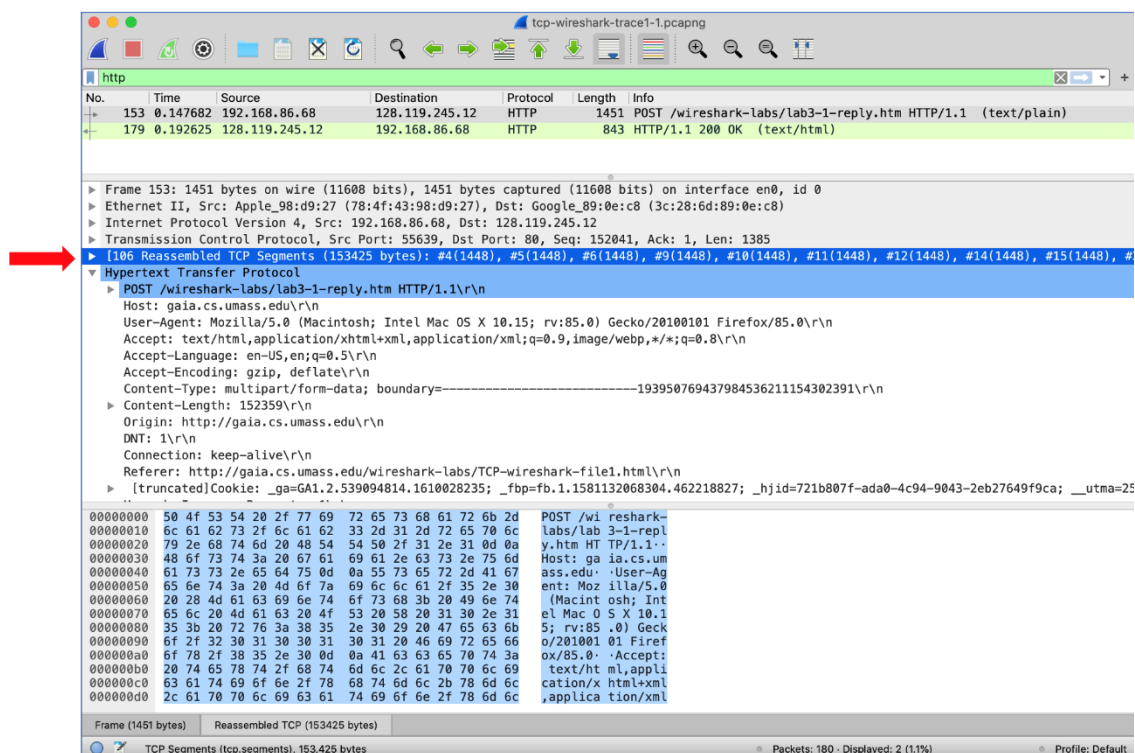


Figure 3: expanding the HTTP POST message that uploaded `alice.txt` from your computer to `gaia.cs.umass.edu`

There are a few things to note here:

- The body of your application-layer HTTP POST message contains the contents of the file `alice.txt`, which is a large file of more than 152K bytes. OK – it’s not *that* large, but it’s going to be too large for this one HTTP POST message to be contained in just one TCP segment!
- In fact, as shown in the Wireshark window in Figure 3 we see that the HTTP POST message was spread across 106 TCP segments. This is shown where the red arrow is placed in Figure 3 [Aside: Wireshark doesn’t have a red arrow like that; we added it to the figure to be helpful ☺]. If you look even more carefully there, you can see that Wireshark is being really helpful to you as well, telling you that the first TCP segment containing the beginning of the POST message is packet #4 in the particular trace for the example in Figure 3, which is the trace `tcp-wireshark-trace1-1` noted in footnote 2. The second TCP segment containing the POST message in packet #5 in the trace, and so on.

Let’s now “get our hands dirty” by looking at some TCP segments.

- First, filter the packets displayed in the Wireshark window by entering “tcp” (lowercase, no quotes, and press return after entering) into the display filter specification window towards the top of the Wireshark window. Your Wireshark display should look something like Figure 4. In Figure 4, we’ve noted the TCP segment that has its SYN bit set – this is the first TCP message in the three-way handshake that sets up the TCP connection to `gaia.cs.umass.edu` that will eventually carry the HTTP POST message and

Continued

the text file. We've also noted the SYNACK segment (the second step in TCP three-way handshake), as well as the TCP segment (packet #4, as discussed above) that carries the POST message and the beginning of the alice.txt file. Of course, the packet numbers will be different in your own trace file, but you should see similar behaviour to that shown in Figures 3 and 4.

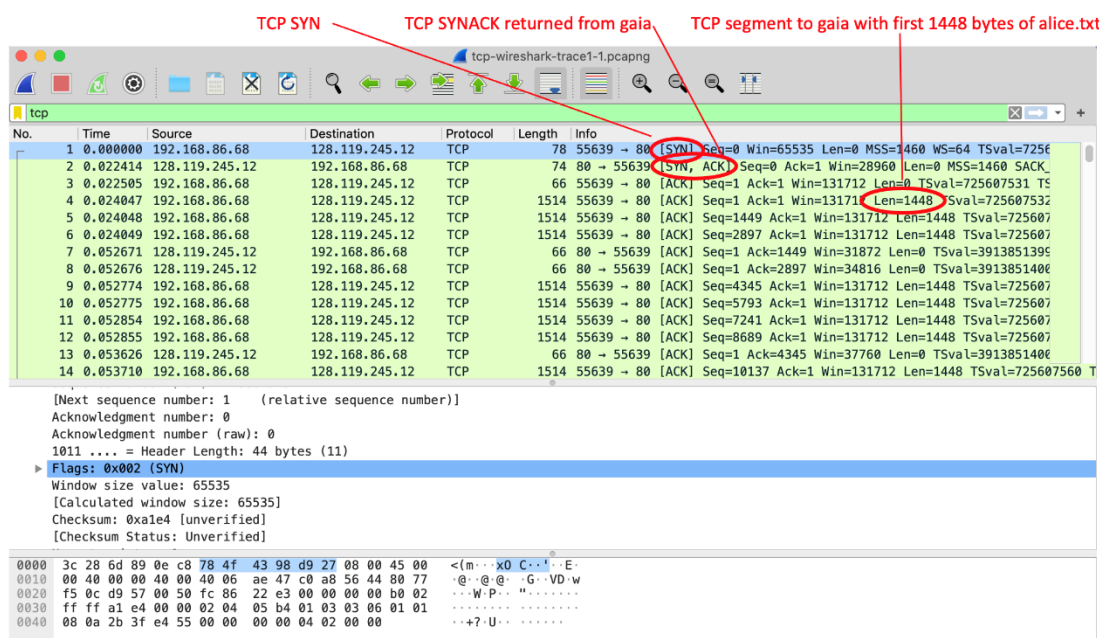


Figure 4: TCP segments involved in sending the HTTP POST message (including the file alice.txt) to gaia.cs.umass.edu

Answer the following questions. Parts that are labelled '[AM]' will be automatically marked and parts that are labelled '[MM]' will be manually marked, while parts with no suffix will not be marked. Where applicable, for manually marked questions, include a screenshot of the packet(s) within the trace that you used to answer a question, or annotate the packet text to explain your answers.

1. What is the IP address and TCP port number used by the client computer (source) that is transferring the alice.txt file to gaia.cs.umass.edu? [AM]
To answer this question, it's probably easiest to select an HTTP message and explore the details of the TCP packet used to carry this HTTP message, using the "details of the selected packet header window" (refer to Figure 2 in the "Getting Started with Wireshark" Lab if you're uncertain about the Wireshark windows).
2. What is the IP address of gaia.cs.umass.edu? On what port number is it sending and receiving TCP segments for this connection? [AM]

Since this lab is about TCP rather than HTTP, now change Wireshark's "listing of captured packets" window so that it shows information about the TCP segments containing the HTTP messages, rather than about the HTTP messages, as in Figure 4 above. (Select Analyze > Enabled

Continued

Protocols, uncheck the HTTP box and select OK.) This is what we're looking for—a series of TCP segments sent between your computer and `gaia.cs.umass.edu`!

3. TCP Basics

Answer the following questions for the TCP segments:

3. What is the *sequence number* of the TCP SYN segment that is used to initiate the TCP connection between the client computer and `gaia.cs.umass.edu`? [AM]
 Note: this is the “raw” sequence number carried in the TCP segment itself; it is NOT the packet # in the “No.” column in the Wireshark window. Remember there is no such thing as a “packet number” in TCP or UDP; as you know, there are sequence numbers in TCP and that’s what we’re after here. Also note that this is not the relative sequence number with respect to the starting sequence number of this TCP session.).
 What is it in this TCP segment that identifies the segment as a SYN segment? [MM]
4. What is the *sequence number* of the SYNACK segment sent by `gaia.cs.umass.edu` to the client computer in reply to the SYN? [AM]
 What is it in the segment that identifies the segment as a SYNACK segment? [MM]
 What is the value of the Acknowledgement field in the SYNACK segment? [AM]
 How did `gaia.cs.umass.edu` determine that value? [MM]
5. What is the sequence number of the TCP segment containing the header of the HTTP POST command? [AM]
 Note that in order to find the POST message header, you’ll need to dig into the packet content field at the bottom of the Wireshark window, *looking for a segment with the ASCII text “POST” within its DATA field*^{5,6}.
6. Consider the TCP segment containing the HTTP “POST” as the first segment in the data transfer part of the TCP connection.
 - At what time was the first segment (the one containing the HTTP POST) in the data transfer part of the TCP connection sent? [AM]
 - At what time was the ACK for this first data-containing segment received? [AM]
 - What is the RTT for this first data-containing segment? [AM]
 - What is the RTT value for the second data-carrying TCP segment and its ACK. [AM]
 - What is the `EstimatedRTT` value (see Section 3.5.3, page 242 in text) after the ACK for the second data-carrying segment is received? [AM]
 Assume that in making this calculation after receiving the ACK for the second segment that the initial value of the `EstimatedRTT` is equal to the measured RTT

⁵ *Hint:* this TCP segment is sent by the client soon (but not always immediately) after the SYNACK segment is received from the server.

⁶ Note that if you filter to only show “http” messages, you’ll see that the TCP segment that Wireshark associates with the HTTP POST message is the *last* TCP segment in the connection (which contains the text at the *end* of `alice.txt`: “THE END”) and *not* the first data-carrying segment in the connection. Students (and teachers!) often find this unexpected and/or confusing.

for the first segment, and then is computed using the `EstimatedRTT` equation on page 242, and a value of $\alpha = 0.125$.

Note: Wireshark has a nice feature that allows you to plot the RTT for each of the TCP segments sent. Select a TCP segment in the “listing of captured packets” window that is being sent from the client to the `gaia.cs.umass.edu` server. Then select: *Statistics->TCP Stream Graph->Round Trip Time Graph*.

7. What is the length (header plus payload) of each of the first four data-carrying TCP segments?⁷ [AM]
8. What is the minimum amount of available buffer space advertised to the client by `gaia.cs.umass.edu` among these first four data-carrying TCP segments?⁸ [AM]
Does the lack of receiver buffer space ever throttle the sender for these first four data-carrying segments? [MM]
9. Are there any retransmitted segments in the trace file? What did you check for (in the trace) in order to answer this question? [MM]

4. TCP congestion control in action [Optional]

Let's now examine the amount of data sent per unit time from the client to the server. Rather than (tediously!) calculating this from the raw data in the Wireshark window, we'll use one of Wireshark's TCP graphing utilities—*Time-Sequence-Graph(Stevens)*—to plot out data.

Select a client-sent TCP segment in the Wireshark's “listing of captured-packets” window corresponding to the transfer of `alice.txt` from the client to `gaia.cs.umass.edu`. Then select the menu: *Statistics->TCP Stream Graph->Time-Sequence-Graph(Stevens)*⁹. You should see a plot that looks similar to the plot in Figure 5, which was created from the captured packets in the packet trace `tcp-wireshark-trace1-1` (see footnote 2). You may have to expand, shrink, and fiddle around with the intervals shown in the axes in order to get your graph to look like Figure 5. (In practice, you're unlikely to see such neat results.)

⁷ The TCP segments in the `tcp-wireshark-trace1-1` trace file are all less than 1480 bytes. This is because the computer on which the trace was gathered has an interface card that limits the length of the maximum IP datagram to 1500 bytes, and there is a *minimum* of 40 bytes of TCP/IP header data. This 1500-byte value is a fairly typical maximum length for an Internet IP datagram.

⁸ Give the Wireshark-reported value for “Window Size Value” which must then be multiplied by the Window Scaling Factor to give the actual number of buffer bytes available at `gaia.cs.umass.edu` for this connection.

⁹ William Stevens wrote the “bible” book on TCP, known as [TCP Illustrated](#).




Figure 5: A sequence-number-versus-time plot (Stevens format) of TCP segments.

Here, each dot represents a TCP segment sent, plotting the sequence number of the segment versus the time at which it was sent. Note that a set of dots stacked above each other represents a series of packets (sometimes called a “fleet” of packets) that were sent back-to-back by the sender.

Use the *Time-Sequence-Graph(Stevens)* plotting tool to view the sequence number versus time plot of segments being sent from the client to the `gaia.cs.umass.edu` server. Can you identify where TCP’s slowstart phase begins and ends, and where congestion avoidance takes over? Comment on ways in which the measured data differs from the idealized behaviour of TCP that we’ve studied in the text.

Note: the reason why this exercise is optional is because your own trace will not nearly be as neat as the one in Fig. 5. That’s not your bad. Can you think of a reason why yours is so different?

Continued





Figure 6: Another view of the same data as in Figure 5.

Continued

Wireshark Lab:

IP v8.1

Supplement to *Computer Networking: A Top-Down Approach*, 8th ed., J.F. Kurose and K.W. Ross



"Tell me and I forget. Show me and I remember. Involve me and I understand." Chinese proverb

© 2005-2021, J.F Kurose and K.W. Ross, All Rights Reserved

In this lab, we'll investigate the IP protocol, focusing on the IPv4 and IPv6 datagram. This lab has three parts. In the first part, we'll analyse packets in a trace of IPv4 datagrams sent and received by the `traceroute` program (the `traceroute` program itself is explored in more detail in the Wireshark ICMP lab). We'll study IP fragmentation in Part 2 of this lab, and take a quick look at IPv6 in Part 3 of this lab.

Before getting started, you'll probably want to review sections 1.4.3 in the text¹⁰ and section 3.4 of [RFC 2151](#) to update yourself on the operation of the `traceroute` program. You'll also want to read Section 4.3 in the text, and probably also have [RFC 791](#) on hand as well, for a discussion of the IP protocol. Although we've removed the topic of IP fragmentation from the 8th edition of our textbook (to make room for new material), you can find material on IP fragmentation from the 7th edition of our textbook (and earlier) at http://gaia.cs.umass.edu/kurose_ross/Kurose_Ross_7th_edition_section_4.3.2.pdf.

¹⁰ References to figures and sections are for the 8th edition of our text, *Computer Networks, A Top-down Approach*, 8th ed., J.F. Kurose and K.W. Ross, Addison-Wesley/Pearson, 2020. Our website for this book is http://gaia.cs.umass.edu/kurose_ross You'll find lots of interesting open material there.

Continued

Capturing packets from an execution of traceroute

In order to generate a trace of IPv4 datagrams for the first two parts of this lab, we'll use the `traceroute` program to send datagrams of two different sizes to `gaia.cs.umass.edu`. Recall that `traceroute` operates by first sending one or more datagrams with the time-to-live (TTL) field in the IP header set to 1; it then sends a series of one or more datagrams towards the same destination with a TTL value of 2; it then sends a series of datagrams towards the same destination with a TTL value of 3; and so on. Recall that a router must decrement the TTL in each received datagram by 1 (actually, RFC 791 says that the router must decrement the TTL by *at least* one). If the TTL reaches 0, the router returns an ICMP message (type 11 – TTL-exceeded) to the sending host. As a result of this behaviour, a datagram with a TTL of 1 (sent by the host executing `traceroute`) will cause the router one hop away from the sender to send an ICMP TTL-exceeded message back to the sender; the datagram sent with a TTL of 2 will cause the router two hops away to send an ICMP message back to the sender; the datagram sent with a TTL of 3 will cause the router three hops away to send an ICMP message back to the sender; and so on. In this manner, the host executing `traceroute` can learn the IP addresses of the routers between itself and the destination by looking at the source IP addresses in the datagrams containing the ICMP TTL-exceeded messages.

Let's run `traceroute` and have it send datagrams of two different sizes. The larger of the two datagram lengths will require `traceroute` messages to be fragmented across multiple IPv4 datagrams.

- Linux/MacOS.** With the Linux/MacOS `traceroute` command, the size of the UDP datagram sent towards the final destination can be explicitly set by indicating the number of bytes in the datagram; this value is entered in the `traceroute` command line immediately after the name or address of the destination. For example, to send `traceroute` datagrams of 2000 bytes towards `gaia.cs.umass.edu`, the command would be:


```
%traceroute gaia.cs.umass.edu 2000
```
- Windows.** The `tracert` program provided with Windows does not allow one to change the size of the ICMP message sent by `tracert`. So it won't be possible to use a Windows machine to generate ICMP messages that are large enough to force IP fragmentation. However, you can use `tracert` to generate small, fixed length packets to perform Part 1 of this lab. At the DOS command prompt enter:


```
>tracert gaia.cs.umass.edu
```

Continued

If you want to do the second part of this lab, you can download a packet trace file that was captured on one of the author's computers¹¹.

Do the following:

- Start up Wireshark and begin packet capture. (*Capture->Start* or click on the blue shark fin button in the top left of the Wireshark window).
- Enter two `tracert` commands, using `gaia.cs.umass.edu` as the destination, the first with a length of 56 bytes. Once that command has finished executing, enter a second `tracert` command for the same destination, but with a length of 3000 bytes.
- Stop Wireshark tracing.

If you're unable to run Wireshark on a live network connection, you can use the packet trace file, *ip-wireshark-trace1-1.pcapng*, referenced in footnote 2. You may well find it valuable to download this trace even if you've captured your own trace and use it, as well as your own trace, as you explore the questions below.

Part 1: Basic IPv4

In your trace, you should be able to see the series of UDP segments (in the case of MacOS/Linux) or ICMP Echo Request messages (Windows) sent by `tracert` on your computer, and the ICMP TTL-exceeded messages returned to your computer by the intermediate routers. In the questions below, we'll assume you're using a MacOS/Linux computer; the corresponding questions for the case of a Windows machine should be clear. Your screen should look similar to the screenshot in Figure 2, where we have used the display filter "`udp | icmp`" (see the light-green-filled display-filter field in Figure 2) so that only UDP and/or ICMP protocol packets are displayed.

¹¹ You can download the zip file <http://gaia.cs.umass.edu/wireshark-labs/wireshark-traces-8.1.zip> and extract the trace file *ip-wireshark-trace1-1.pcapng*. This trace file can be used to answer these Wireshark lab questions without actually capturing packets on your own. The trace was made using Wireshark running on one of the author's computers, while performing the steps in this Wireshark lab. Once you've downloaded a trace file, you can load it into Wireshark and view the trace using the *File* pull down menu, choosing *Open*, and then selecting the trace file name.

Continued

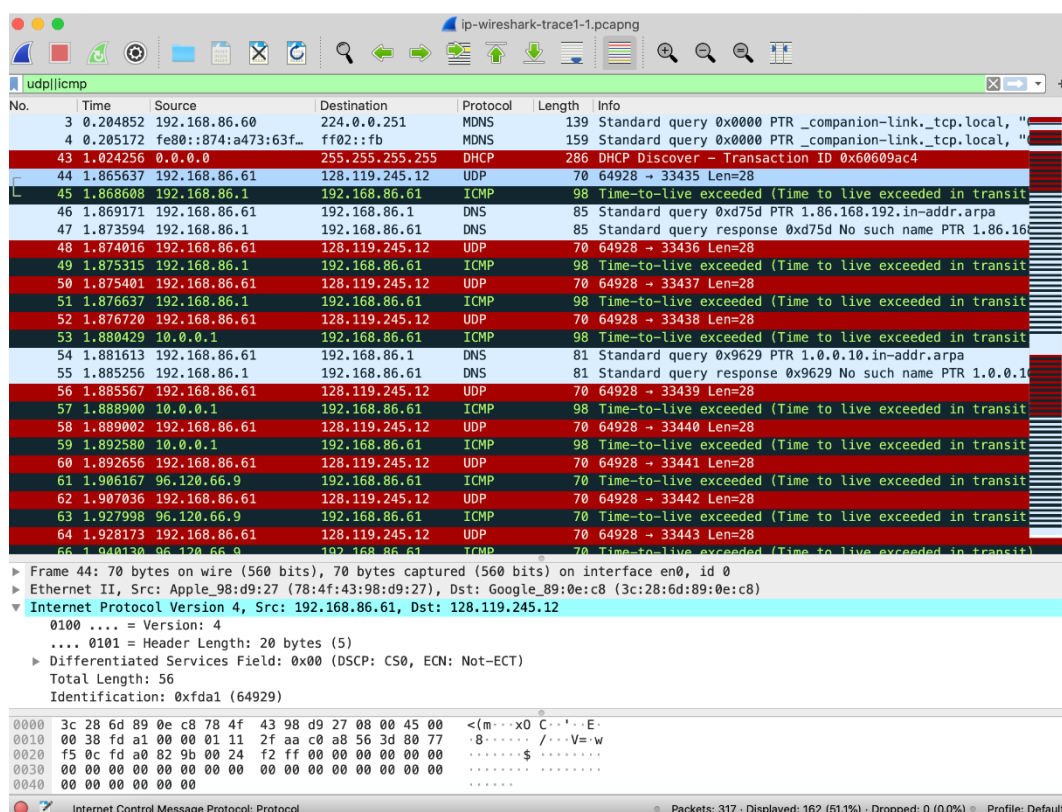


Figure 2: Wireshark screenshot, showing UDP and ICMP packets in the tracefile *ip-wireshark-trace1-1.pcapng*

Answer the following questions.

10. Select the first UDP segment sent by your computer via the `traceroute` command to `gaia.cs.umass.edu` and expand the Internet Protocol part of the packet in the packet details window. What is the IP address of your computer? [AM]
11. What is the value in the time-to-live (TTL) field in this IPv4 datagram's header? [AM]
12. What is the value in the upper layer protocol field in this IPv4 datagram's header? [Note: the answers for Linux/macOS differ from Windows here]. [AM]
13. How many bytes are in the IP header? How many bytes are in the payload of the IP datagram? [AM]
14. Has this IP datagram been fragmented? [AM]
15. Explain how you determined whether or not the datagram has been fragmented. [MM]

Next, let's look at the *sequence* of UDP segments being sent from your computer via `traceroute`, destined to `128.119.245.12`. The display filter that you can enter to do this is "`ip.src==192.168.86.61 and ip.dst==128.119.245.12 and udp and icmp`". This will allow you to easily move sequentially through just the datagrams containing just these segments. Your screen should look similar to Figure 3.

Continued

No.	Time	Source	Destination	Protocol	Length	Info
44	1.865637	192.168.86.61	128.119.245.12	UDP	70	64928 → 33435 Len=28
48	1.874016	192.168.86.61	128.119.245.12	UDP	70	64928 → 33436 Len=28
50	1.875401	192.168.86.61	128.119.245.12	UDP	70	64928 → 33437 Len=28
52	1.876720	192.168.86.61	128.119.245.12	UDP	70	64928 → 33438 Len=28
56	1.885567	192.168.86.61	128.119.245.12	UDP	70	64928 → 33439 Len=28
58	1.889002	192.168.86.61	128.119.245.12	UDP	70	64928 → 33440 Len=28
60	1.892656	192.168.86.61	128.119.245.12	UDP	70	64928 → 33441 Len=28
62	1.907036	192.168.86.61	128.119.245.12	UDP	70	64928 → 33442 Len=28
64	1.928173	192.168.86.61	128.119.245.12	UDP	70	64928 → 33443 Len=28
67	1.940279	192.168.86.61	128.119.245.12	UDP	70	64928 → 33444 Len=28
69	1.951481	192.168.86.61	128.119.245.12	UDP	70	64928 → 33445 Len=28
71	1.965335	192.168.86.61	128.119.245.12	UDP	70	64928 → 33446 Len=28
73	1.975799	192.168.86.61	128.119.245.12	UDP	70	64928 → 33447 Len=28
75	1.991739	192.168.86.61	128.119.245.12	UDP	70	64928 → 33448 Len=28
77	2.008887	192.168.86.61	128.119.245.12	UDP	70	64928 → 33449 Len=28
79	2.025877	192.168.86.61	128.119.245.12	UDP	70	64928 → 33450 Len=28

▼ Internet Protocol Version 4, Src: 192.168.86.61, Dst: 128.119.245.12

- 0100 = Version: 4
- 0101 = Header Length: 20 bytes (5)
- Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
- Total Length: 56
- Identification: 0xfdb6 (64950)
- Flags: 0x0000
- Fragment offset: 0
- Time to live: 8
- Protocol: UDP (17)
- Header checksum: 0x2895 [validation disabled]
- [Header checksum status: Unverified]
- Source: 192.168.86.61
- Destination: 128.119.245.12

▼ User Datagram Protocol, Src Port: 64928, Dst Port: 33456

- Source Port: 64928
- Destination Port: 33456

Figure 3: Wireshark screenshot, showing up segments in the tracefile *ip-wireshark-trace1-1.pcapng* using the display filter `ip.src==192.168.86.61` and `ip.dst==128.119.245.12` and `udp` and `!icmp`

- Which fields in the IP datagram *always* change from one datagram to the next within this series of UDP segments sent by your computer destined to 128.119.245.12, via `traceroute`? [MM]
- Which fields in this sequence of IP datagrams (containing UDP segments) stay constant? Why? [MM]

Continued

18. Describe the pattern you see in the values in the Identification field of the IP datagrams being sent by your computer. [MM]

Now let's take a look at the ICMP packets being returned to your computer by the intervening routers where the TTL value was decremented to zero (and hence caused the ICMP error message to be returned to your computer). The display filter that you can use to show just these packets is "ip.dst==192.168.86.61 and icmp".

19. What is the upper layer protocol specified in the IP datagrams returned from the routers? [Note: the answers for Linux/macOS differ from Windows here]. [MM]
20. Are the values in the Identification fields (across the sequence of all of ICMP packets from all of the routers) similar in behaviour to your answer to question 18 above? [MM]
21. Are the values of the TTL fields similar, across all of ICMP packets from all of the routers? [MM]

Network performance during loadshedding

There is a widespread suspicion and anecdotal experience that the networks are slower during loadshedding, especially during the higher stages of loadshedding. This may be regarding reducing bandwidth, stability of the connection, cell coverage or fibre cable switches upstream having no on-line UPS but only a generator that takes time to start, and more.

22. Devise an experiment that will falsify, or validate the hypothesis "The network is slower during higher stages of loadshedding (stage > 4) than lower ones (stage < 4)". That is, write down a procedure about what needs to be done to investigate this to obtain a conclusion 'yes this does happen' or 'no, the performance is the same'. This procedure should be written in such a way that it can be given to a classmate and they'd be able to carry out the experiment you devised.¹² [MM]
23. Did your network have problems when uploading the Gutenberg file? Try the filters from the table below and report whether there was anything out of the ordinary. For instance, applying the tcp.analysis.flags may show a "TCP Dup ACK", a duplicate ACK received, tcp.analysis.ack_rtt > 0.1 may show packets as well (look at the bottom-right for number of packets and percentage of slow RTT packets). [MM]

¹² You are encouraged to try it out as well. But, since the stages are predictably unreliable, it's not possible to make it part of the assignment.

Selection of filters that can be applied to your trace, which may indicate network problems.¹³

Name	Filter	Description
Connection Reset	<code>tcp.flags.reset==1</code>	Should never happen mid conversation. Usually a network failure or timeout.
Bad TCP packets	<code>tcp.analysis.flags && !tcp.analysis.window_update</code>	Sliding windows update is normal, anything else shouldn't be there.
Slow RTT	<code>tcp.analysis.ack_rtt > 0.1</code>	If this starts high or climbs throughout trace it's a sign of congestion.
TCP WAITING FOR APP	<code>tcp.time_delta > 0.1</code>	Is the App (client or server) so busy it can't respond. Suggests an application problem
SLOW DNS	<code>dns.time > 1</code>	Slow name resolution. Generally impacts every command.

THE END

¹³ Table from <https://portal.perforce.com/s/article/2956> and more suggestions can be found there, among many places.